

# class 10 - candy

jack olmstead

```
url <- "https://raw.githubusercontent.com/fivethirtyeight/data/master/candy-power-ranking/"  
candy <- read.csv(url, row.names = 1)
```

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

For some cursed reason, the tasteless, unwashed FifeThirtyEight analysts who put this frivolous exercise together didn't include Hot Tamales in the dataset. Probably because its win-rate was so game-breaking.

```
candy["Haribo Sour Bears",]$winpercent
```

```
[1] 51.41243
```

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat",]$winpercent
```

```
[1] 76.7686
```

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars",]$winpercent
```

```
[1] 49.6535
```

```
library(skimr)
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	
None	

#### Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

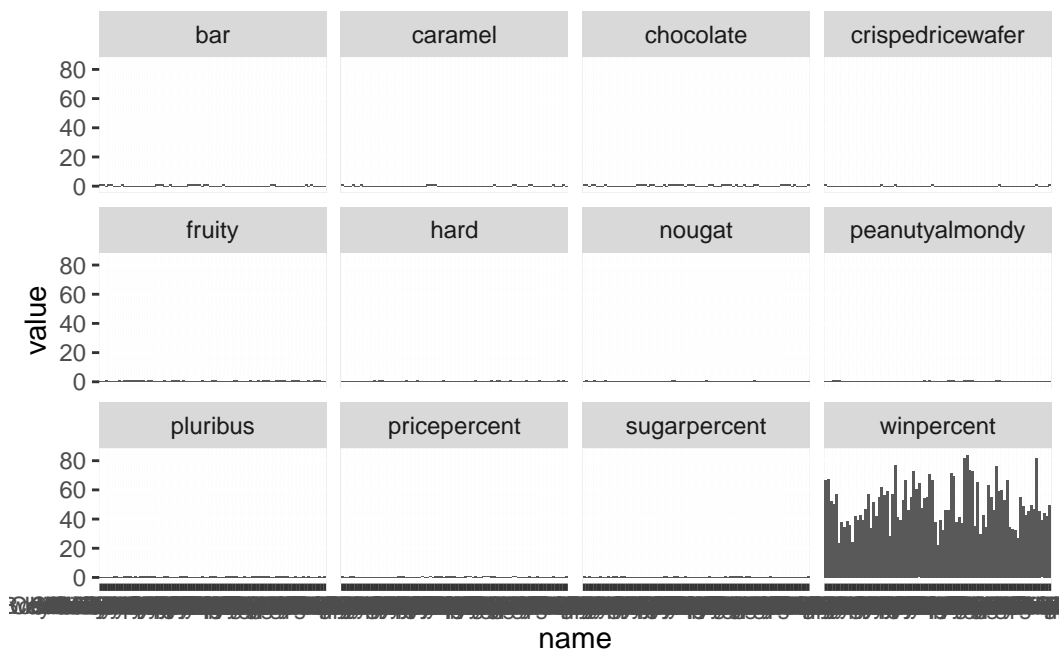
Win percent definite has a different scale than everything else.

```
library(tidyverse)
```

```
-- Attaching packages ----- tidyverse 1.3.2 --
v ggplot2 3.4.0      v purrr   1.0.1
v tibble  3.1.8      v dplyr   1.1.0
v tidyr   1.3.0      v stringr 1.5.0
v readr   2.1.3      v forcats 1.0.0
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
```

```
candy$name <- rownames(candy)
neat <- gather(candy, "var", "value", -name)
```

```
ggplot(neat) +
  aes(x=name, y=value) +
  geom_col() +
  facet_wrap(~var)
```




Q7. What do you think a zero and one represent for the candy\$chocolate column?

- 1: the candy is chocolate.
- 0: the candy is not chocolate.

Q8. Plot a histogram of winpercent values

```
ggplot(candy) +  
  aes(x=winpercent) +  
  geom_histogram()
```

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



Q9. Is the distribution of winpercent values symmetrical?

Looks to be long-tailed on the right side.

Q10. Is the center of the distribution above or below 50%?

```
median(candy$winpercent)
```

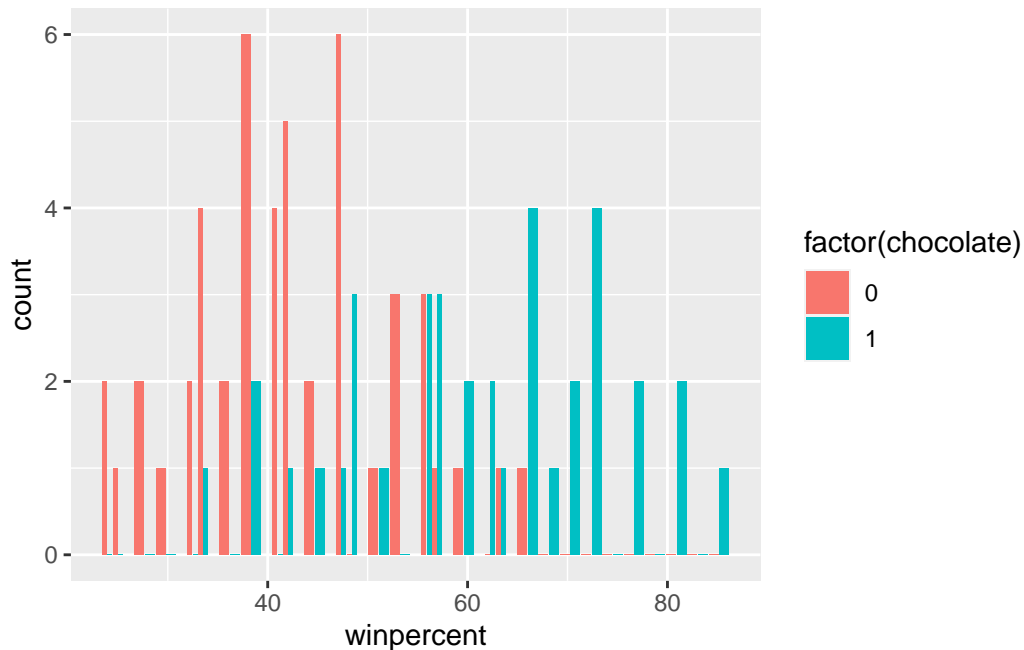
```
[1] 47.82975
```

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Definitely a chocolate bias

```
ggplot(candy) +  
  aes(x=winpercent, fill=factor(chocolate)) +  
  geom_histogram(position="dodge2")
```

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



Q.12. Is this difference statistically significant?

Very.

```
t.test(  
  candy[candy$chocolate == 1,]$winpercent,  
  candy[candy$chocolate == 0,]$winpercent  
)
```

Welch Two Sample t-test

data: candy[candy\$chocolate == 1,]\$winpercent and candy[candy\$chocolate == 0,]\$winpercent

```
t = 7.3031, df = 67.539, p-value = 4.164e-10
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 13.64744 23.91110
sample estimates:
mean of x mean of y
 60.92153  42.14226
```

Q13. What are the five least liked candy types in this set?

```
sorter <- sort(candy$winpercent, index.return=1, decreasing = T)
candy.sorted <- candy[sorter$ix,]
tail(candy.sorted)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Root Beer Barrels	0	0	0		0	0
Jawbusters	0	1	0		0	0
Super Bubble	0	1	0		0	0
Chiclets	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Nik L Nip	0	1	0		0	0

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Root Beer Barrels				0	1	0	1	0.732		0.069
Jawbusters				0	1	0	1	0.093		0.511
Super Bubble				0	0	0	0	0.162		0.116
Chiclets				0	0	0	1	0.046		0.325
Boston Baked Beans				0	0	0	1	0.313		0.511
Nik L Nip				0	0	0	1	0.197		0.976

	winpercent	name
Root Beer Barrels	29.70369	Root Beer Barrels
Jawbusters	28.12744	Jawbusters
Super Bubble	27.30386	Super Bubble
Chiclets	24.52499	Chiclets
Boston Baked Beans	23.41782	Boston Baked Beans
Nik L Nip	22.44534	Nik L Nip

Q14. What are the top 5 all time favorite candy types out of this set?

```
head(candy.sorted)
```

```
chocolate fruity caramel peanut almond nougat
```

Reese's Peanut Butter cup	1	0	0	1	0
Reese's Miniatures	1	0	0	1	0
Twix	1	0	1	0	0
Kit Kat	1	0	0	0	0
Snickers	1	0	1	1	1
Reese's pieces	1	0	0	1	0


	crisp	edrice	wafer	hard	bar	pluribus	sugar	percent
Reese's Peanut Butter cup		0	0	0		0		0.720
Reese's Miniatures		0	0	0		0		0.034
Twix		1	0	1		0		0.546
Kit Kat		1	0	1		0		0.313
Snickers		0	0	1		0		0.546
Reese's pieces		0	0	0		1		0.406

	price	percent	winpercent	name
Reese's Peanut Butter cup	0.651	84.18029		Reese's Peanut Butter cup
Reese's Miniatures	0.279	81.86626		Reese's Miniatures
Twix	0.906	81.64291		Twix
Kit Kat	0.511	76.76860		Kit Kat
Snickers	0.651	76.67378		Snickers
Reese's pieces	0.651	73.43499		Reese's pieces

Q15. Make a first barplot of candy ranking based on winpercent values.

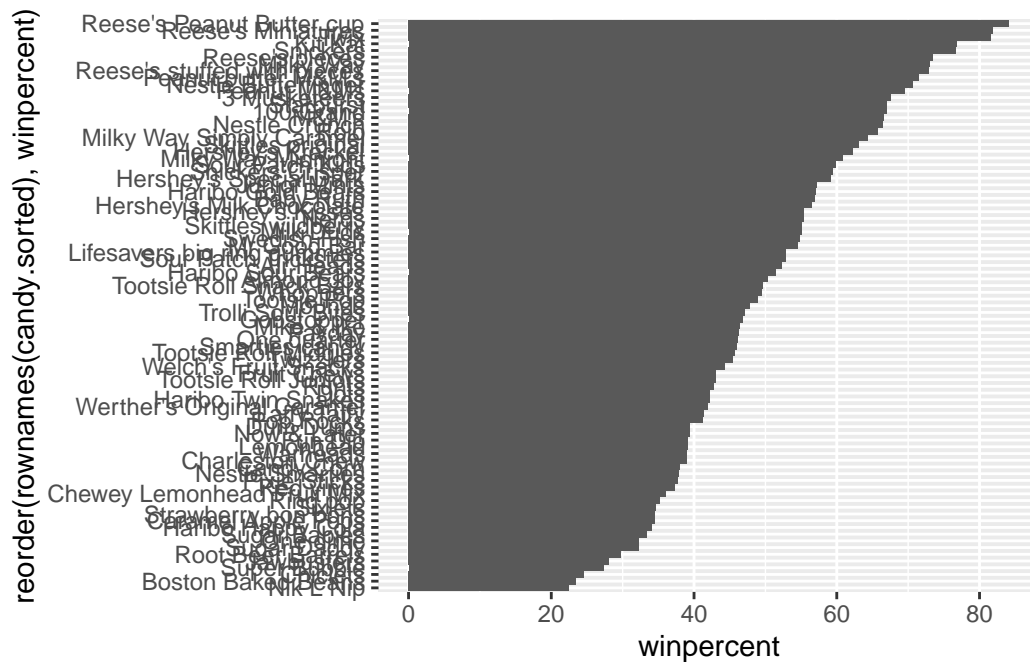
```
ggplot(candy.sorted) +
  aes(x=winpercent, y=rownames(candy.sorted)) +
  geom_col()
```



Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by winpercent?

```
ggplot(candy.sorted) +  
  aes(x=winpercent, y=reorder(rownames(candy.sorted), winpercent)) +  
  geom_col()
```






Let's make some useful color vectors!

```
my.colors <- rep("black", nrow(candy))
my.colors[as.logical(candy$chocolate)] = "chocolate"
my.colors[as.logical(candy$bar)] = "brown"
my.colors[as.logical(candy$fruity)] = "pink"
```

```
ggplot(candy.sorted) +
  aes(x=winpercent, y=reorder(rownames(candy.sorted), winpercent)) +
  geom_col(fill=my.colors)
```



Q17. What is the worst ranked chocolate candy?

Nik L Nip, whatever that is.

Q18. What is the best ranked fruity candy?

Snickers, apparently? Although I don't think of this as a fruity candy.

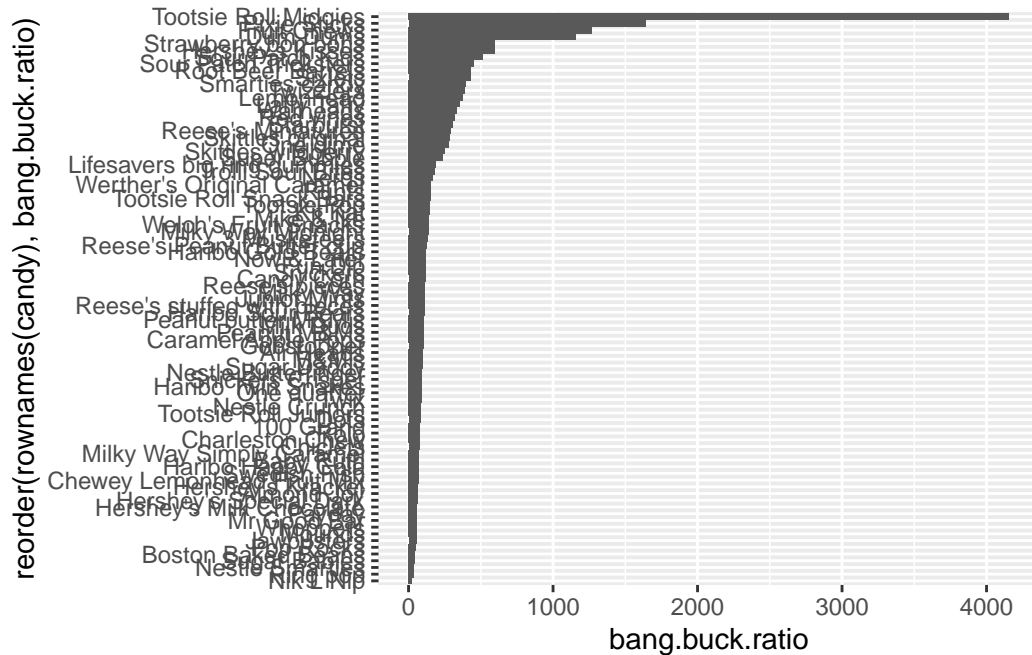
```
library("ggrepel")

ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my.colors) +
  geom_text_repel(col=my.colors, size=3.3, max.overlaps = 5)
```

Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

```
candy$bang.buck.ratio <- candy$winpercent / candy$pricepercent

ggplot(candy) +
  aes(x=bang.buck.ratio, y=reorder(rownames(candy), bang.buck.ratio)) +
  geom_col()
```



Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
candy.sort.indicies.price <- sort(candy$pricepercent, decreasing = T, index.return=T)
head(candy[candy.sort.indicies.price$ix, ])
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Nestle Smarties	1	0	0		0	0
Ring pop	0	1	0		0	0
Hershey's Krackel	1	0	0		0	0
Hershey's Milk Chocolate	1	0	0		0	0
Hershey's Special Dark	1	0	0		0	0

	crispedrice	wafer	hard bar	pluribus	sugarpercent	
Nik L Nip		0	0	0	1	0.197
Nestle Smarties		0	0	0	1	0.267
Ring pop		0	1	0	0	0.732
Hershey's Krackel		1	0	1	0	0.430
Hershey's Milk Chocolate		0	0	1	0	0.430
Hershey's Special Dark		0	0	1	0	0.430

	pricepercent	winpercent	name
Nik L Nip	0.976	22.44534	Nik L Nip


Nestle Smarties	0.976	37.88719	Nestle Smarties
Ring pop	0.965	35.29076	Ring pop
Hershey's Krackel	0.918	62.28448	Hershey's Krackel
Hershey's Milk Chocolate	0.918	56.49050	Hershey's Milk Chocolate
Hershey's Special Dark	0.918	59.23612	Hershey's Special Dark
	bang.buck.ratio		
Nik L Nip	22.99728		
Nestle Smarties	38.81884		
Ring pop	36.57073		
Hershey's Krackel	67.84802		
Hershey's Milk Chocolate	61.53649		
Hershey's Special Dark	64.52737		

### exploring the dimensionality of these data

```
library(corrplot)
```

corrplot 0.92 loaded

```
cij <- cor(candy[, -13])
corrplot(cij)
```



```
pcs <- prcomp(candy[, -13], scale=T)
summary(pcs)
```

Importance of components:


	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0938	1.2127	1.13054	1.0787	0.98027	0.93656	0.81530
Proportion of Variance	0.3372	0.1131	0.09832	0.0895	0.07392	0.06747	0.05113
Cumulative Proportion	0.3372	0.4503	0.54866	0.6382	0.71208	0.77956	0.83069

	PC8	PC9	PC10	PC11	PC12	PC13
Standard deviation	0.78462	0.68466	0.66328	0.57829	0.43128	0.39534
Proportion of Variance	0.04736	0.03606	0.03384	0.02572	0.01431	0.01202
Cumulative Proportion	0.87804	0.91410	0.94794	0.97367	0.98798	1.00000

Plot PCs


```
plot(pcs$x[,1], pcs$x[,2], col=my.colors, pch=16)
```



Add these PC coordinates to our candy data.


```
candy <- cbind(candy, pcs$x[, 1:3])
```

```
p <- ggplot(candy) +  
  aes(x=PC1, y=PC2,  
       size=winpercent/100,  
       color=my.colors) +  
  geom_point()  
p
```




Last, let's look at the loadings of our stuffs.

```
par(mar=c(8,4,2,2))
barplot(pcs$rotation[,1], las=2, ylab="PC1 Contribution")
```



```
barplot(pcs$rotation[,2], las=2, ylab="PC2 Contribution")
```





```
# library(patchwork)
# p1+p2
```

Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity, hard, pluribus, and bang-buck ratio positively drive PC1 scores. Everything else negatively drives PC1 scores. Interesting, there aren't any variables that are *not* correlated with this PC.